

Demonstrating Clustering using K-Means Algorithm

Azhar Chowdhury

Statistical Programmer|Data Scientist[SAS, Python, R]

The following meta summary is helpful to follow the problem at hand and it's solution.

Dataset The Portugal Electricity Consumption dataset

Domain Time Series

ML Algorithm K-Means Clustering

Components Cluster. Silhouette Score

Problem Description Perform clustering using K-Means algorithm, Identify the number of clusters for a given dataset

```
[1]: import os
      os.environ["OMP_NUM_THREADS"] = "1"
```

The above code will prevent from over-consuming CPU or memory, or when we are getting OpenMP-related warnings, setting `OMP_NUM_THREADS = 1` as a safe baseline - as I am running PY programs on Windows 11.

```
[2]: pathToFile = r'C://Users//azhar//Downloads//'
      fileName = 'LD2011_2014.txt'

      import numpy as np
      import pandas as pd
      from sklearn.cluster import KMeans
      import matplotlib.pyplot as plt
      import random
      from sklearn.metrics import silhouette_score
      from sklearn.cluster import AgglomerativeClustering
      random.seed(42)
```

```
[3]: # Replace "," by ".", otherwise the numbers will be in the form 2,3445 instead
      ↪ of 2.3445
      import fileinput

      with fileinput.FileInput(pathToFile+fileName, inplace=True, backup='.bak') as
      ↪ file:
          for line in file:
```

```
print(line.replace(",", "."), end='')
```

```
[4]: import pandas as pd
data = pd.read_csv(pathToFile+fileName, sep=";", index_col=0)
```

```
[5]: data.head(2)
data.tail(2)
data.shape
data.info()
data.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 140256 entries, 2011-01-01 00:15:00 to 2015-01-01 00:00:00
Columns: 370 entries, MT_001 to MT_370
dtypes: float64(370)
memory usage: 397.0+ MB
```

```
[5]:
```

	MT_001	MT_002	MT_003	MT_004 \
count	140256.000000	140256.000000	140256.000000	140256.000000
mean	3.970785	20.768480	2.918308	82.184490
std	5.983965	13.272415	11.014456	58.248392
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	2.844950	0.000000	36.585366
50%	1.269036	24.893314	1.737619	87.398374
75%	2.538071	29.871977	1.737619	115.853659
max	48.223350	115.220484	151.172893	321.138211

	MT_005	MT_006	MT_007	MT_008 \
count	140256.000000	140256.000000	140256.000000	140256.000000
mean	37.240309	141.227385	4.521338	191.401476
std	26.461327	98.439984	6.485684	121.981187
min	0.000000	0.000000	0.000000	0.000000
25%	15.853659	71.428571	0.565291	111.111111
50%	39.024390	157.738095	2.826456	222.222222
75%	54.878049	205.357143	4.522329	279.461279
max	150.000000	535.714286	44.657999	552.188552

	MT_009	MT_010 ...	MT_361	MT_362 \
count	140256.000000	140256.000000 ...	140256.000000	140256.000000
mean	39.975354	42.205152 ...	218.213701	37607.987537
std	29.814595	33.401251 ...	204.833532	38691.954832
min	0.000000	0.000000 ...	0.000000	0.000000
25%	13.986014	9.677419 ...	5.710207	0.000000
50%	40.209790	40.860215 ...	131.334761	24100.000000
75%	57.692308	61.290323 ...	403.283369	54800.000000
max	157.342657	198.924731 ...	852.962170	192800.000000


	MT_363	MT_364	MT_365	MT_366 \
--	--------	--------	--------	----------

count	140256.000000	140256.000000	140256.000000	140256.000000
mean	1887.427366	2940.031734	65.413150	9.269709
std	1801.486488	2732.251967	65.007818	10.016782
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	13.037810	0.000000
50%	1050.632911	2136.363636	31.290743	7.021650
75%	3312.236287	5363.636364	108.213820	11.702750
max	7751.054852	12386.363636	335.071708	60.269163

	MT_367	MT_368	MT_369	MT_370
count	140256.000000	140256.000000	140256.000000	140256.000000
mean	424.262904	94.704717	625.251734	8722.355145
std	274.337122	80.297301	380.656042	9195.155777
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	30.050083	83.944282	0.000000
50%	525.899912	76.794658	758.064516	0.000000
75%	627.743635	151.919866	875.366569	17783.783784
max	1138.718174	362.270451	1549.120235	30918.918919

[8 rows x 370 columns]

```
[6]: data_example = data.loc['2012-01-01 00:15:00':'2012-01-03 00:00:
    ↳00'][['MT_001', 'MT_002']]
data_example.plot()
plt.show()
```



```
[7]: data2011 = data.loc['2011-01-01 00:15:00':'2012-01-01 00:00:00']
data2012 = data.loc['2012-01-01 00:15:00':'2013-01-01 00:00:00']
data2013 = data.loc['2013-01-01 00:15:00':'2014-01-01 00:00:00']
data2014 = data.loc['2014-01-01 00:15:00':'2015-01-01 00:00:00']
```

```
[8]: # Check number of days
print(data2011.shape[0]/96)
print(data2012.shape[0]/96)
print(data2013.shape[0]/96)
print(data2014.shape[0]/96)
```

```
365.0
366.0
365.0
365.0
```

```
[9]: # See number of clients with 0 demand per year
print(sum(data2011.mean()==0))
print(sum(data2012.mean()==0))
print(sum(data2013.mean()==0))
print(sum(data2014.mean()==0))
```

```
210
```

37
21
1

```
[10]: import pandas as pd
clients = data2011.columns
clients_no_demand = clients[data2013.mean()==0] # clients with 0 demand
#data_13_14 = data2013.concat(data2014) # appending 2013 and 2014
data_13_14 = pd.concat([data2013, data2014], axis=0)
data_13_14 = data_13_14.drop(clients_no_demand, axis=1) # drop clients with 0 demand
    ↪demand
print(data_13_14.shape)
print(sum(data_13_14.mean()==0)) # check that there are no clients with 0 demand
```

(70080, 349)
0

```
[11]: data = data_13_14.copy() # weekdays weekends, data2011, data2012, data2013, ↪
    ↪data2014
data['hour'] = data.index.map(lambda x: x[11:])
data.head(5)
```

```
[11]:
```

	MT_001	MT_002	MT_003	MT_004	MT_005	\
2013-01-01 00:15:00	2.538071	22.759602	2.606429	138.211382	63.414634	
2013-01-01 00:30:00	1.269036	22.759602	2.606429	138.211382	63.414634	
2013-01-01 00:45:00	2.538071	22.759602	2.606429	134.146341	60.975610	
2013-01-01 01:00:00	1.269036	23.470839	2.606429	130.081301	56.097561	
2013-01-01 01:15:00	3.807107	23.470839	2.606429	130.081301	58.536585	

	MT_006	MT_007	MT_008	MT_009	MT_010	\
2013-01-01 00:15:00	255.952381	4.522329	239.057239	57.692308	78.494624	
2013-01-01 00:30:00	264.880952	5.652911	228.956229	57.692308	76.344086	
2013-01-01 00:45:00	250.000000	5.652911	239.057239	54.195804	76.344086	
2013-01-01 01:00:00	226.190476	6.218202	249.158249	50.699301	75.268817	
2013-01-01 01:15:00	229.166667	6.783493	239.057239	57.692308	74.193548	

	...	MT_362	MT_363	MT_364	MT_365	\
2013-01-01 00:15:00	...	22300.0	886.075949	1000.000000	16.949153	
2013-01-01 00:30:00	...	21000.0	864.978903	909.090909	18.252934	
2013-01-01 00:45:00	...	18200.0	860.759494	840.909091	16.949153	
2013-01-01 01:00:00	...	15800.0	860.759494	840.909091	16.949153	
2013-01-01 01:15:00	...	15000.0	793.248945	818.181818	16.949153	

	MT_366	MT_367	MT_368	MT_369	MT_370	\
2013-01-01 00:15:00	6.436513	616.330114	76.794658	731.671554	8086.486486	
2013-01-01 00:30:00	3.510825	564.530290	76.794658	727.272727	8086.486486	
2013-01-01 00:45:00	5.851375	590.869183	68.447412	730.205279	7848.648649	
2013-01-01 01:00:00	4.095963	575.065847	58.430718	722.873900	7848.648649	


```
2013-01-01 01:15:00  4.095963  570.676032  60.100167  748.533724  7610.810811
```

```
             hour
2013-01-01 00:15:00  00:15:00
2013-01-01 00:30:00  00:30:00
2013-01-01 00:45:00  00:45:00
2013-01-01 01:00:00  01:00:00
2013-01-01 01:15:00  01:15:00
```

```
[5 rows x 350 columns]
```

```
[12]: # Getting average curves per client
datagrouped = data.groupby("hour")
average_curves = datagrouped.agg("mean")
average_curves.shape
average_curves_norm = average_curves/(average_curves.mean())

average_curves_norm[['MT_001', 'MT_002', 'MT_369', 'MT_370']].plot()
plt.show()
```




```
[13]: client = 'MT_022'
oneClient = data_13_14[client]
X = [] # a list of arrays, each array being a normalized curve for a day
for J in range(2*365):
    X.extend([np.array(oneClient[J*96:(J+1)*96])])
```

```
[14]: import numpy as np
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
import random
# from sklearn.metrics import silhouette_score
# from sklearn.cluster import AgglomerativeClustering
random.seed(43)

trueK = 11 # Desired number of clusters
dim = 2 # dimension of feature space
#n = 40 # points per cluster
trueCentroids = np.random.rand(trueK, dim) # generating centroids at random
print(trueCentroids)

if dim==2:
    plt.scatter(trueCentroids[:,0],trueCentroids[:,1],)
    plt.xlim(-0.1,1.1)
    plt.ylim(-0.1,1.1)
    plt.show()
```

```
[0.68350152 0.50044166]
[0.32660083 0.79046578]
[0.78053065 0.22799233]
[0.12871697 0.61836057]
[0.17087974 0.66002271]
[0.93675917 0.01834244]
[0.9139862  0.75203116]
[0.28655132 0.75612555]
[0.8641899  0.51231176]
[0.196591   0.88015142]
[0.19226105 0.99728998]]
```



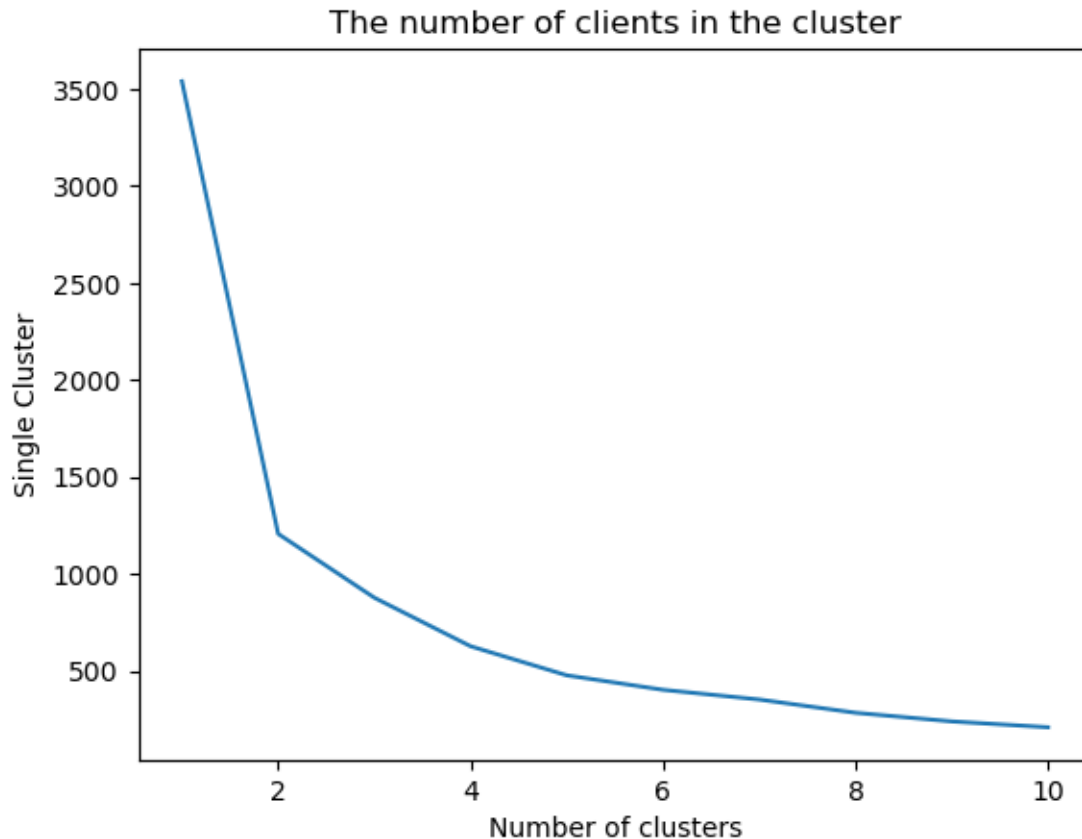
```
[22]: import os
import numpy as np
from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=trueK, n_init=10, random_state=43).fit(X)

print(kmeans.inertia_)
```

179.30226039855134

```
[23]: from sklearn.cluster import KMeans
_clusters = []
for i in range(1, trueK):
    kmeans = KMeans(n_clusters = i, n_init=10, init = 'k-means++', random_state=
↳ 43)
    kmeans.fit(X)
    _clusters.append(kmeans.inertia_)

plt.title("The number of clients in the cluster")
plt.plot(range(1, 11), _clusters,)
plt.xlabel('Number of clusters')
plt.ylabel('Single Cluster')
plt.show()
```





```
[18]: import os
os.environ["OMP_NUM_THREADS"] = "1" # Avoid MKL memory leak warning on Windows

import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans

_clusters = []
trueK = 11 # Make sure this is defined appropriately
for i in range(1, trueK):
    kmeans = KMeans(n_clusters=i, init='k-means++', random_state=43, n_init=10)
    ↪ # Set n_init explicitly
    kmeans.fit(X)
    _clusters.append(kmeans.inertia_)

plt.title("The number of clients in the cluster")
plt.plot(range(1, trueK), _clusters)
plt.xlabel('Number of clusters')
plt.ylabel('Inertia')
plt.show()
```



To determine a number of clusters we guess the number of clusters. In a day of 24 hours, usage may vary. Lets assume that the usage varies from 1 to 11 hours in a day. If we fit the KMeans object with the number of cluster (11) and request for inertia_ (KMeans.inertia_) for each cluster we'll get 11 inertia values. Then if we plot them the graph will reveal an elbow shape. The point where the elbow is determined is our desired number of clusters. In our case it is 2.

```
[19]: from sklearn import metrics
      from sklearn.metrics import pairwise_distances

      X = average_curves_norm.copy()
```

```
[26]: import numpy as np
      from sklearn.cluster import KMeans
      kmeans = KMeans(n_clusters=2, n_init=10, random_state=1).fit(X)
      labels = kmeans.labels_
      metrics.silhouette_score(X, labels, metric='euclidean')
```

```
[26]: 0.5676594904032211
```

A Silhouette score close to 1 is a well spreaded cluster. Since our score is positive and close to 1 it is well spreaded, non-overlapping and observations do not seem to be in a wrong cluster. Comparing

to inertia score, it is clear that I was not able to set the data properly fit to the model and hence the inertia was not calculated correctly. However, if it were calculated correctly Silhouette method is a better evaluator.